

Mohamed Aklamaash M.R

 [Github](#)  [Linkedin](#)  [Mail](#)  [+91 63692 02355](#)  [Portfolio](#)

SUMMARY

AI engineer in the final year of an integrated M.Sc., with a year of production experience at a YC-backed startup, where I took secrets detection (DLP) to 97% recall and 94% F1 and built the analytics that track how every AI agent in an org gets used. The rest of my time goes into LLM tooling I ship in the open, like an adversarial-robustness benchmark on PyPI and a GPT trained from scratch, which is where I dig deeper into models and evaluation. Comfortable in PyTorch and the everyday LLM stack, and about as happy debugging a flaky pipeline as designing the model side. Graduating May 2027.

EXPERIENCE

Unbound Security (YC S24), Software Engineering Intern (Summer)

May 2026 – Jun 2026

Remote / Bengaluru, India

- Added secrets and DLP scanning to the platform with TruffleHog, then tuned it to **97% recall** and **94% F1**. A follow-up false-positive cleanup cut the alert noise security teams had been complaining about.
- Tracked down a classifier field-name mismatch (path vs files vs paths) that was silently skipping policy checks, and closed the coverage gap.
- Shipped BYOC deployment so enterprise customers can run the product inside their own cloud: EKS with Helm and ArgoCD for GitOps, Authentik wired into the auth flow for SSO, and Prometheus plus Grafana for monitoring.

Unbound Security (YC S24), Software Engineering Intern

May 2025 – Dec 2025

Remote / Bengaluru, India

- Pushed the DLP detection pipelines past **90% classification accuracy**, and built AI-usage governance analytics covering every AI agent in the org.
- Migrated the core product off GCP onto AWS across a distributed microservices architecture, cutting infra cost by **up to 20%** with no hit to reliability or performance.
- Rebuilt billing to handle invoice-based, usage-metered pricing on a high-throughput service, and owned it end to end from design through production.

PROJECTS

Redharness: LLM adversarial-robustness benchmark, published on PyPI *Python · LLMs · evaluation*



- A reproducible benchmark for LLM robustness across jailbreaks, prompt injection, and data leakage. Every score ships with a (dataset, judge, metric) triple, so two runs are actually comparable instead of vibes.

RAMPART: Regime-aware multi-agent portfolio allocation *Python · RL · PPO*



- A multi-agent allocator where a supervisor routes capital across specialist strategies under hard risk guardrails. Validated on purged walk-forward backtests, so it stays research and simulation, not trading advice.

Nanocoder-math: Decoder-only GPT trained from scratch *Python · PyTorch · Transformers*



- Trained a decoder-only GPT from scratch for code and math on a single device: a custom 16k BPE tokenizer, a streaming data pipeline, and controlled ablations on optimizer, data, and architecture (MPS and CUDA).

TECHNICAL SKILLS

AI & LLMs: LangChain, LlamaIndex, LangSmith, OpenAI / Anthropic / Gemini / Groq APIs, RAG, Embeddings, Vector search, Semantic search, Knowledge graphs, Prompt engineering, Fine-tuning, Function/tool calling, Structured outputs, LLM evaluation & benchmarking, Guardrails, Red-teaming, Adversarial robustness

Agents & Orchestration: Multi-agent systems, LangGraph, CrewAI, AutoGen, Agno, SmolAgents, Supervisor / planner-executor patterns, Memory & tool use, Low-code automation (n8n, Langflow)

ML & Deep Learning: PyTorch, TensorFlow, Scikit-learn, Transformers & attention, Training from scratch, Tokenization (BPE), Reinforcement learning (PPO, policy-gradient, actor-critic), Supervised & unsupervised learning, Classification, Hyperparameter tuning, Ablation studies, NumPy, Pandas

Languages: Python, C++, Java, Go, TypeScript, SQL

Cloud, Data & MLOps: AWS, Kubernetes (EKS), Docker, Helm, ArgoCD (GitOps), CI/CD, FastAPI / REST APIs, PostgreSQL, MySQL, MongoDB, Redis, Pinecone, Kafka, RabbitMQ, Prometheus, Grafana, Git

EDUCATION

PSG College of Technology, Coimbatore, India

Aug 2022 – May 2027 (Expected)

M.Sc. (Integrated) Data Science · CGPA: 8.18 / 10

Relevant coursework: Machine Learning, Deep Learning, Probability & Statistics, Data Structures & Algorithms, DBMS, Distributed Computing, Operating Systems